

AVOID TRAPS IN YOUR STACK AND YOUR CULTURE

THE MICROLITH

Jason Melo - Chief Architect

**MICROSERVICES AREN'T FREE...AND
WE HAVE YET TO "SOLVE" IT.**

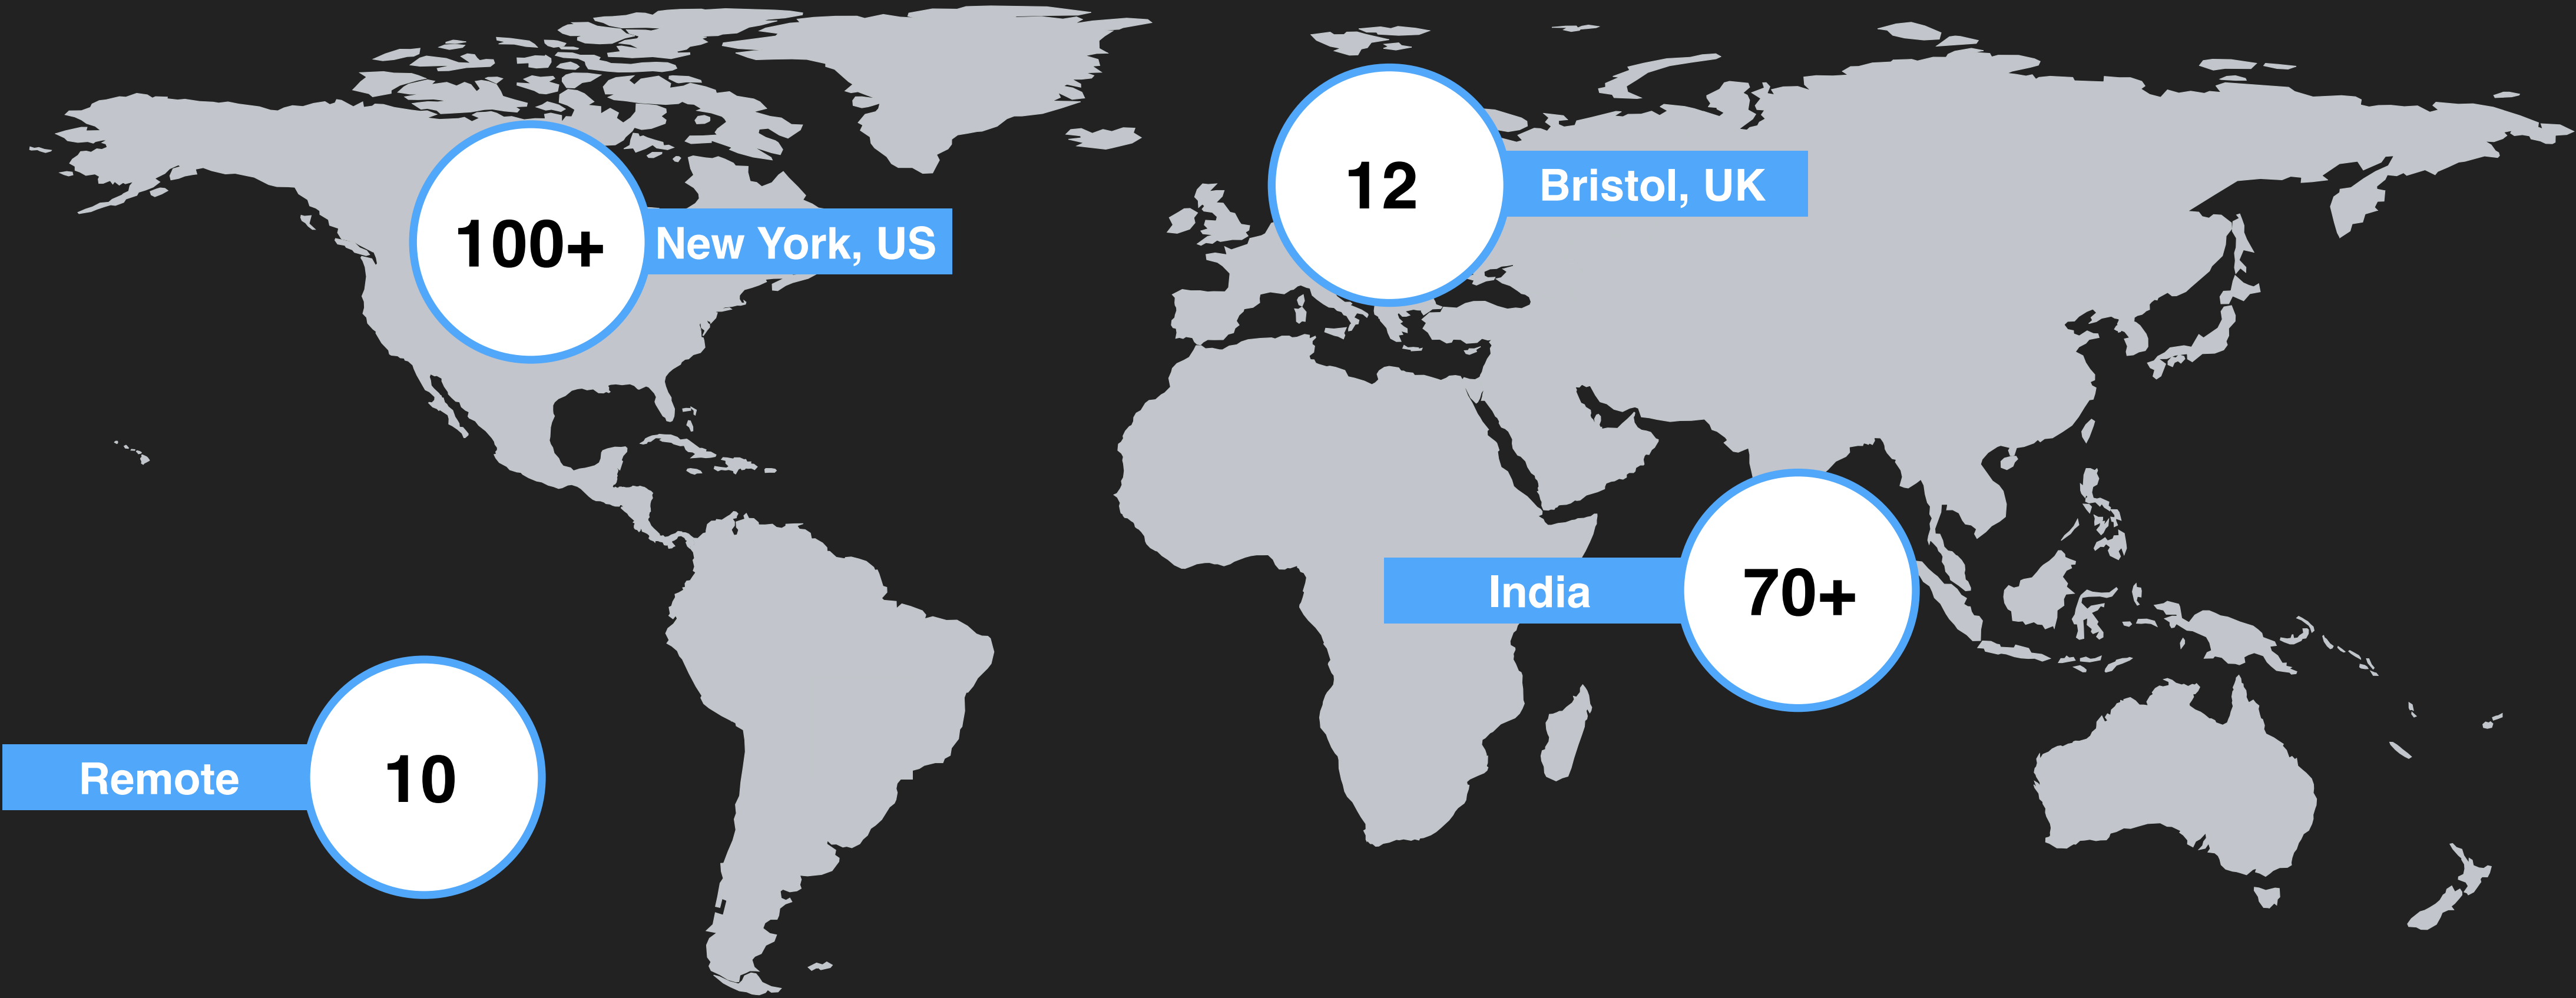
Everyone

Microlith, defined as a Distributed Monolith
...a large number of "micro" services
experiencing coupled computing & data
dependency hell when executing top-level
business transactions

IN CONCLUSION

- ▶ Realize your services as a fully distributed architecture
- ▶ Microservices don't come for free - the investment in tech & culture is not trivial
- ▶ Go faster - your PaaS is a unit of computing, not a pattern to force your developers into

LIFION BY ADP – AROUND THE WORLD



SOME NUMBERS

- ▶ We deal with Humans, their professional & personal interactions from birth to grave on a global scale with localized precision - massive complexity



40m Active Users

The infographic consists of a white circle with a blue border containing the text '40m'. To its right is a blue horizontal rectangle containing the text 'Active Users' in white.



500k Clients

The infographic consists of a white circle with a blue border containing the text '500k'. To its right is a blue horizontal rectangle containing the text 'Clients' in white.



30k Logins/sec

The infographic consists of a white circle with a blue border containing the text '30k'. To its right is a blue horizontal rectangle containing the text 'Logins/sec' in white.



▶ ABOUT 11 MONTHS OLD

▶ 90+ SERVICES,
MOSTLY DOCKER

▶ DISTRIBUTED MODEL

▶ EACH ENVIRONMENT
~ 1000 INSTANCES

L I F
I O N

by ADP

SOME CONTEXT

- ▶ Essentially building a massive distributed JIT compiler & runtime in the cloud. Patterns that we've seen in the industry don't quite fit, some unique challenges are presented
- ▶ At Lifion all services cannot be run in a fully independent way...Metadata is everything - our users create data structures as well as user interfaces via Metadata
- ▶ We now think of our Microservices as a distributed architecture, would you want your Cassandra cluster communicating & distributing data using synchronous http REST?...same rules should apply to your application

THE JOURNEY



PHASE 1 – CULTURAL MIND-SHIFT IS EVERYTHING

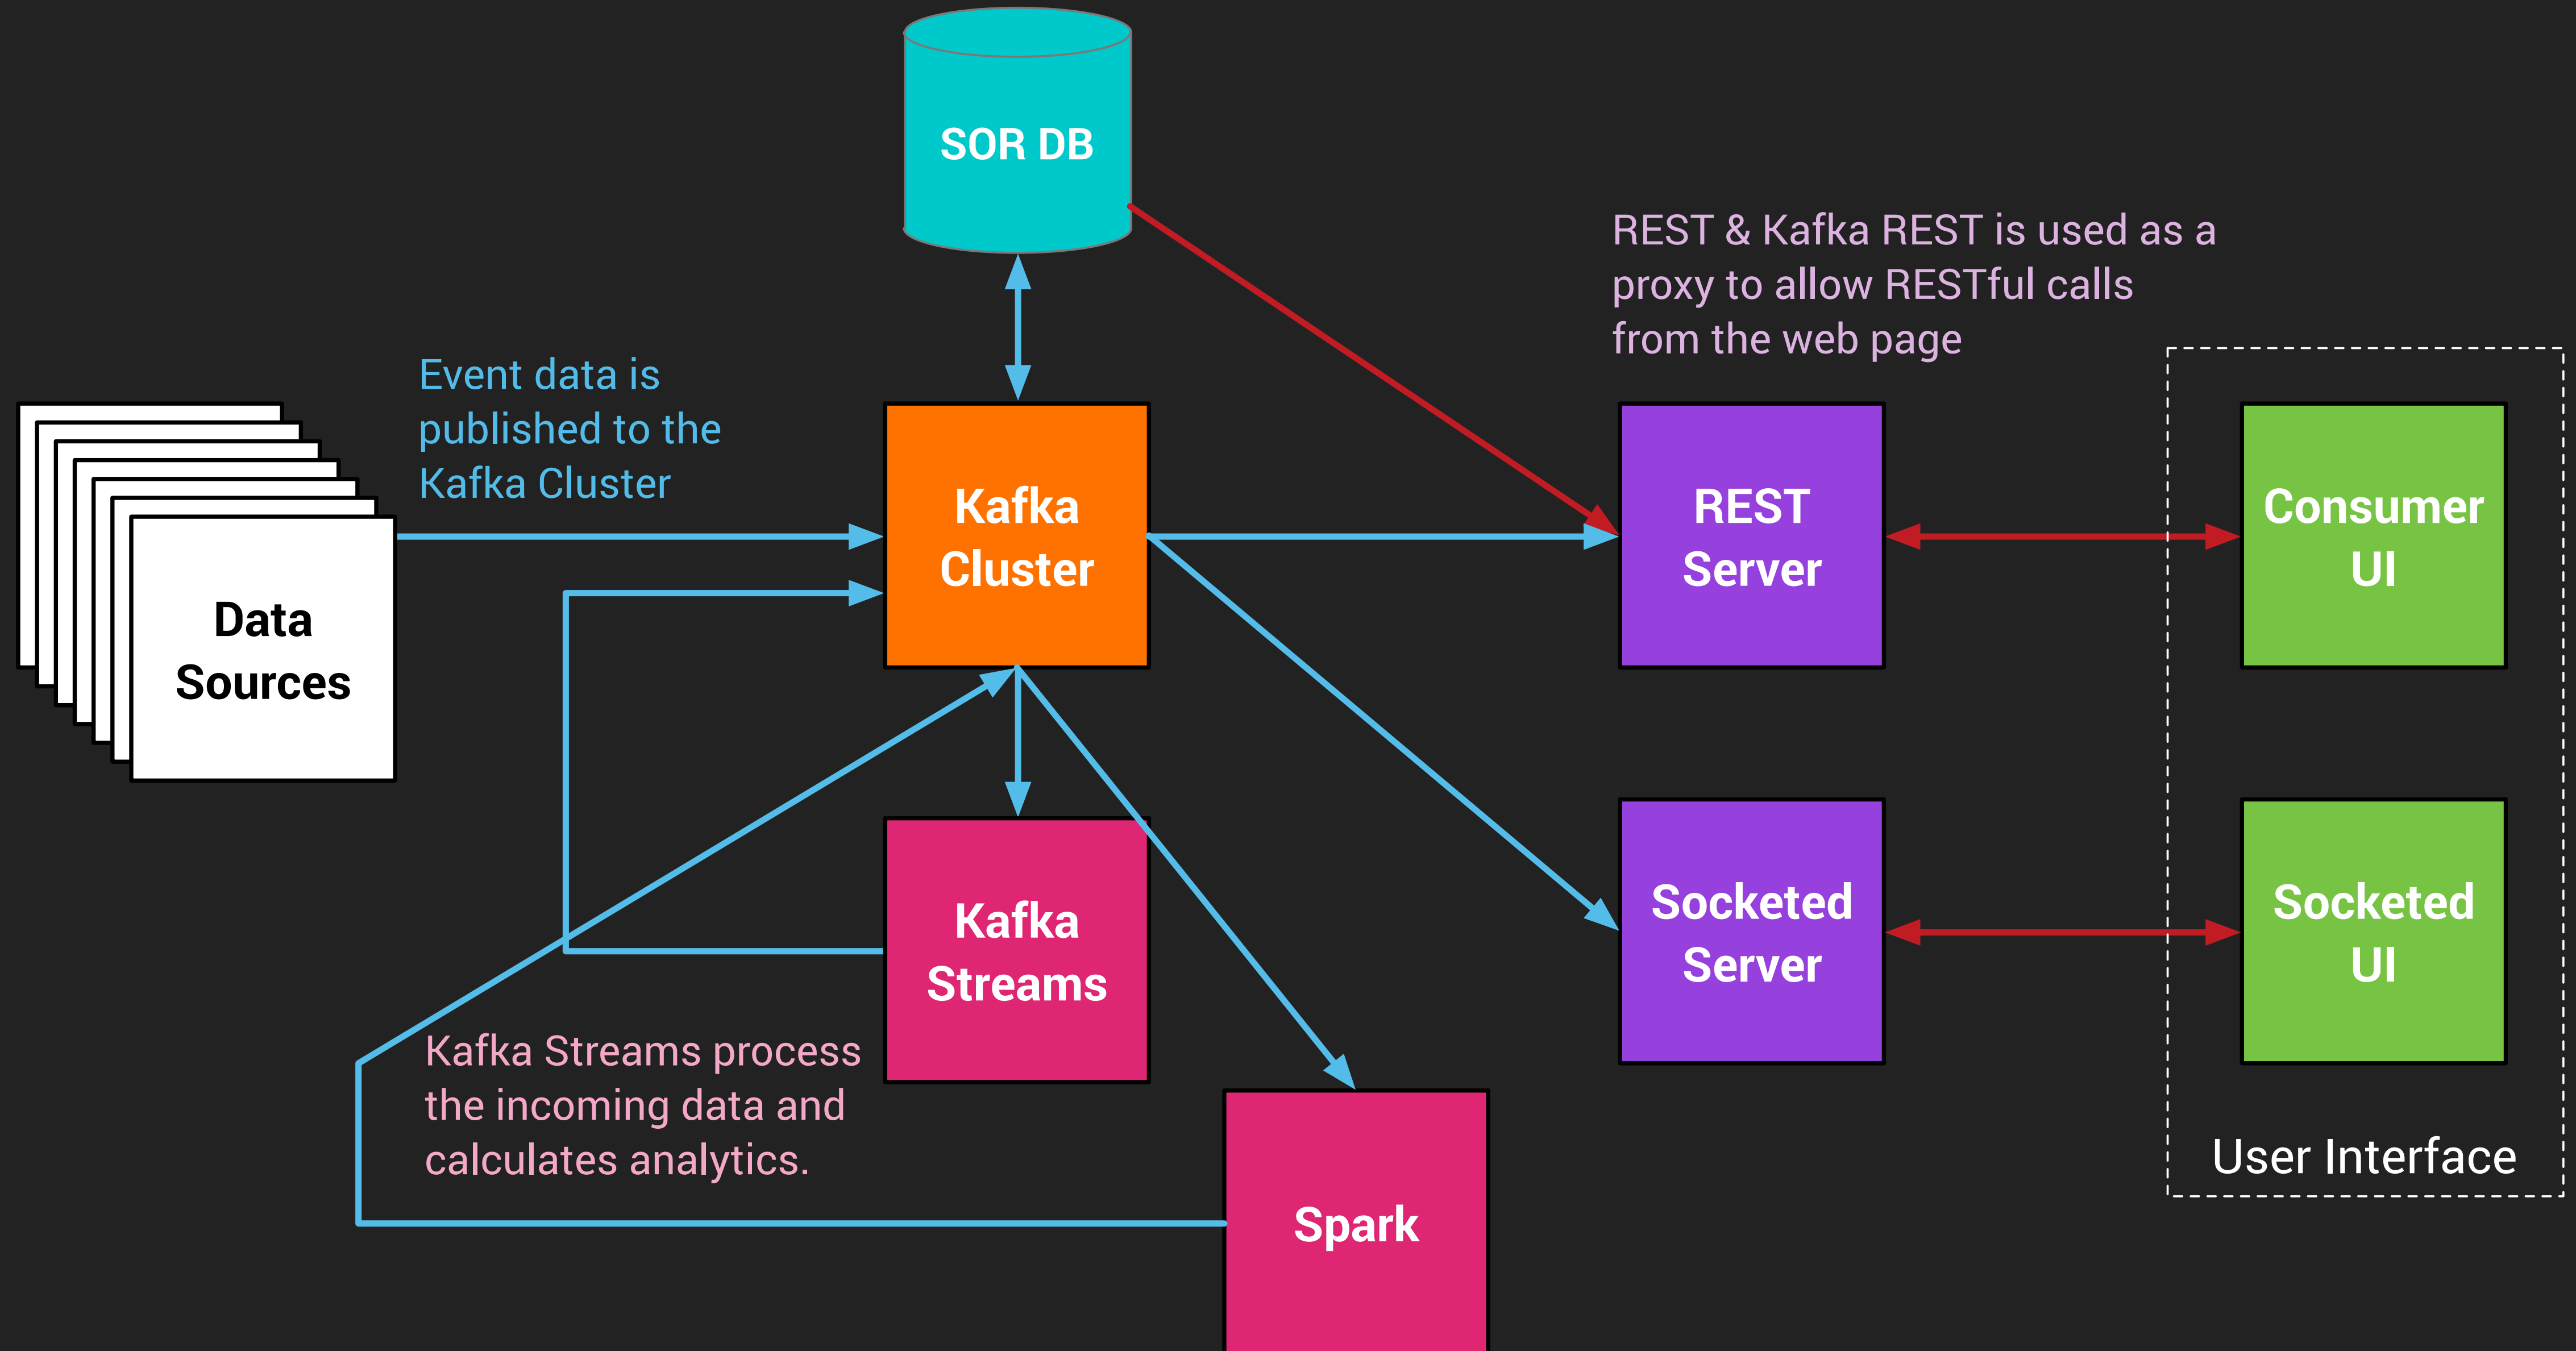
- ▶ Tooling & attitude necessary to go from 1 to 100's of services
- ▶ Gates of service entry
 - ▶ stateful vs stateless services
 - ▶ devops is not a team, its you!
 - ▶ service deltas are additive
 - ▶ well defined service ownership & accountability
- ▶ Embrace Chaos
 - ▶ tons of databases - Kafka, Couchbase, Cassandra, Graph, Time-series, MySQL, Memcache, Mongo
 - ▶ A service team defines their own Architecture & DB needs, we just don't care if they've satisfied our gates
 - ▶ multiple runtimes - Node (multi-ver), JVM (Java, Groovy & Scala) & Python

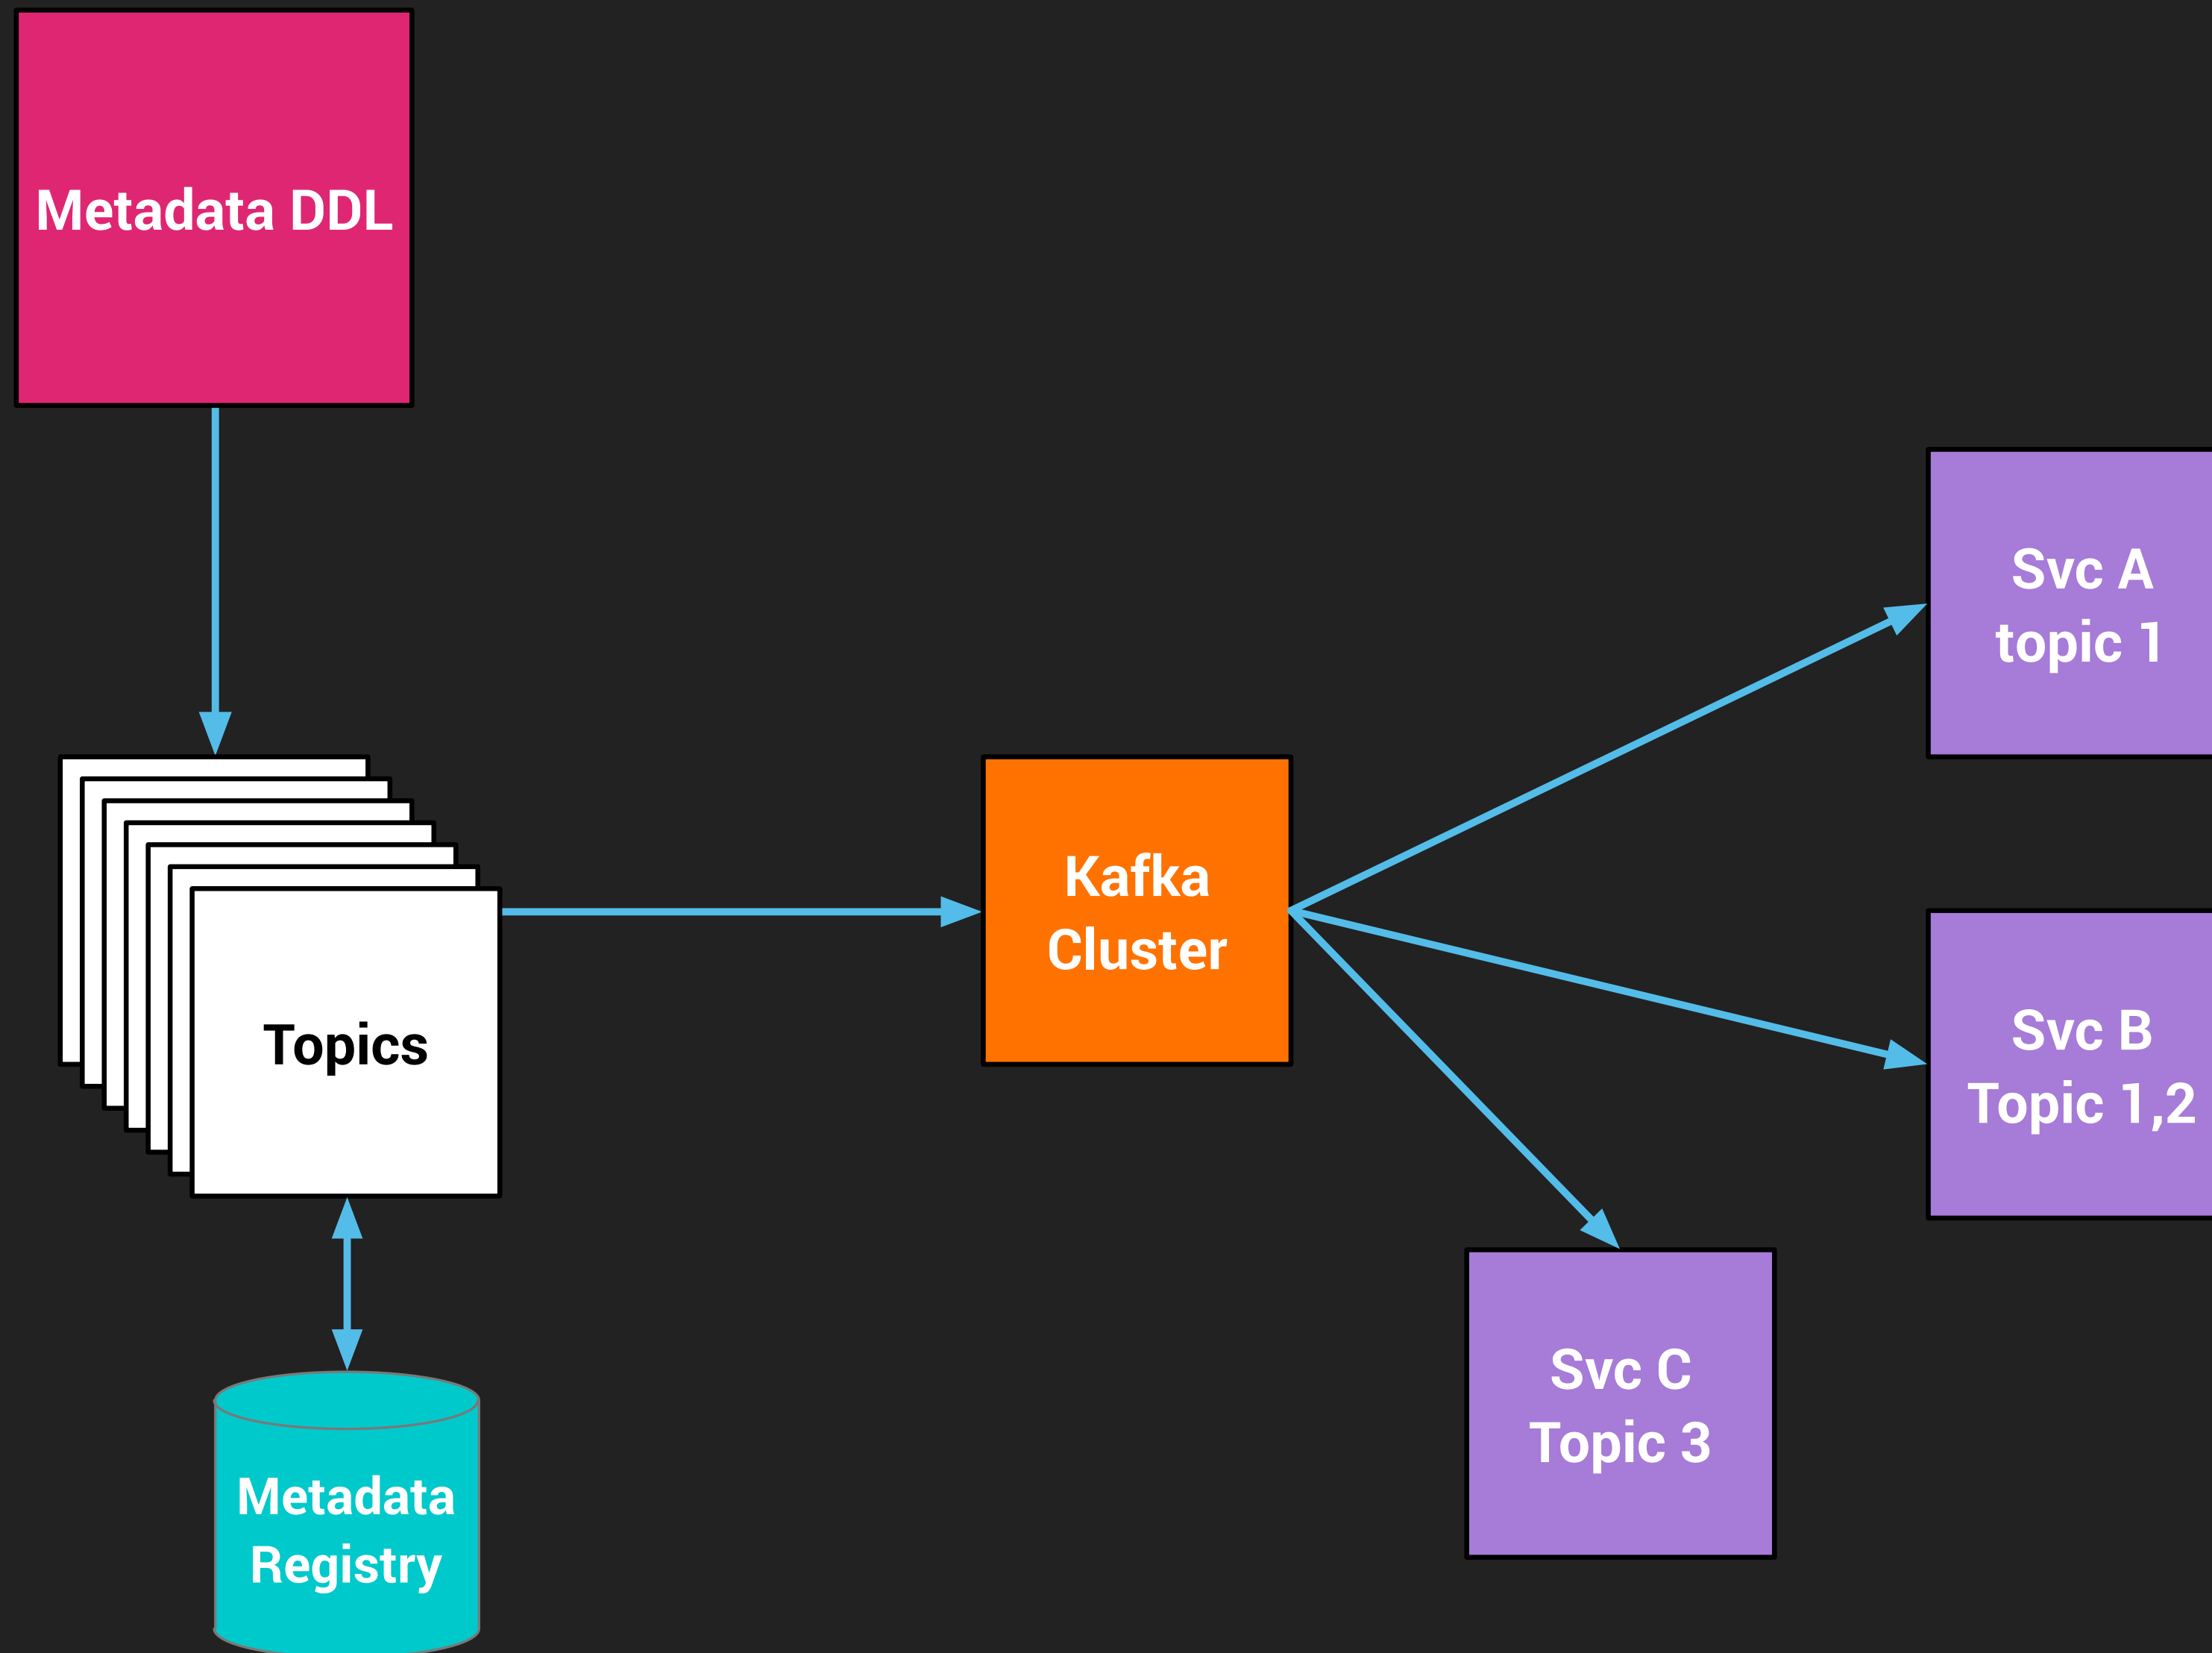
PHASE 1 – STRANGLING THE MONOLITH

- ▶ Teams now control their services with domain boundaries, guard rails & build pipelines
- ▶ Very high network chattiness dependent on synchronous REST
- ▶ Metadata makes mocks almost useless
 - ▶ difficult to develop locally with indicative use cases
 - ▶ a top-level business transaction require synchronous calls across a minimum of 8+ services
- ▶ Service registration & discovery is elegant allowing us to solve the scaling problem & cloud challenges

PHASE 2

With Event messaging patterns over REST, governance, aggregated logging, transaction tracers & architecture supporting streaming data we are decoupling the Microlith - fun, not easy





- ▶ Services register their intent to consume Metadata which are paired with our pipeline Topics
- ▶ As Metadata DDL occurs the distributed in-process cache in the SDK is invalidated (SDK's only for Metadata, not as clients)
- ▶ Significantly reduces synchronous chatter & eliminates the need to consume Metadata over the wire in realtime

LOGGING & ANALYTICS

- ▶ ELK, Kafka, Spark & S3 - aggregated logging solution
 - ▶ standardized data pipeline in use for compute, orchestration & logging
 - ▶ extremely high message throughput for realtime log ingestion & analysis
 - ▶ realtime security & performance analytics
- ▶ Standard logging libraries & formats governed across all service owners
- ▶ Realtime log analysis on deployment in blue/green scenarios using the same analytics platform

SCHEDULING & ORCHESTRATION

- ▶ Jenkins, Docker UCP/Swarm on AWS, Cloud Formation & Ansible
 - ▶ known, reliable and un-complicated
 - ▶ you need to work hard at getting persistence & scaling right
- ▶ Distributed data streaming, messaging & analytics pipeline runs on the above solidly, However:
 - ▶ Beta program with DC/OS, Triton & Kubernetes
 - ▶ known, a bit unstable (improving rapidly), higher operational complexity
 - ▶ local development challenging

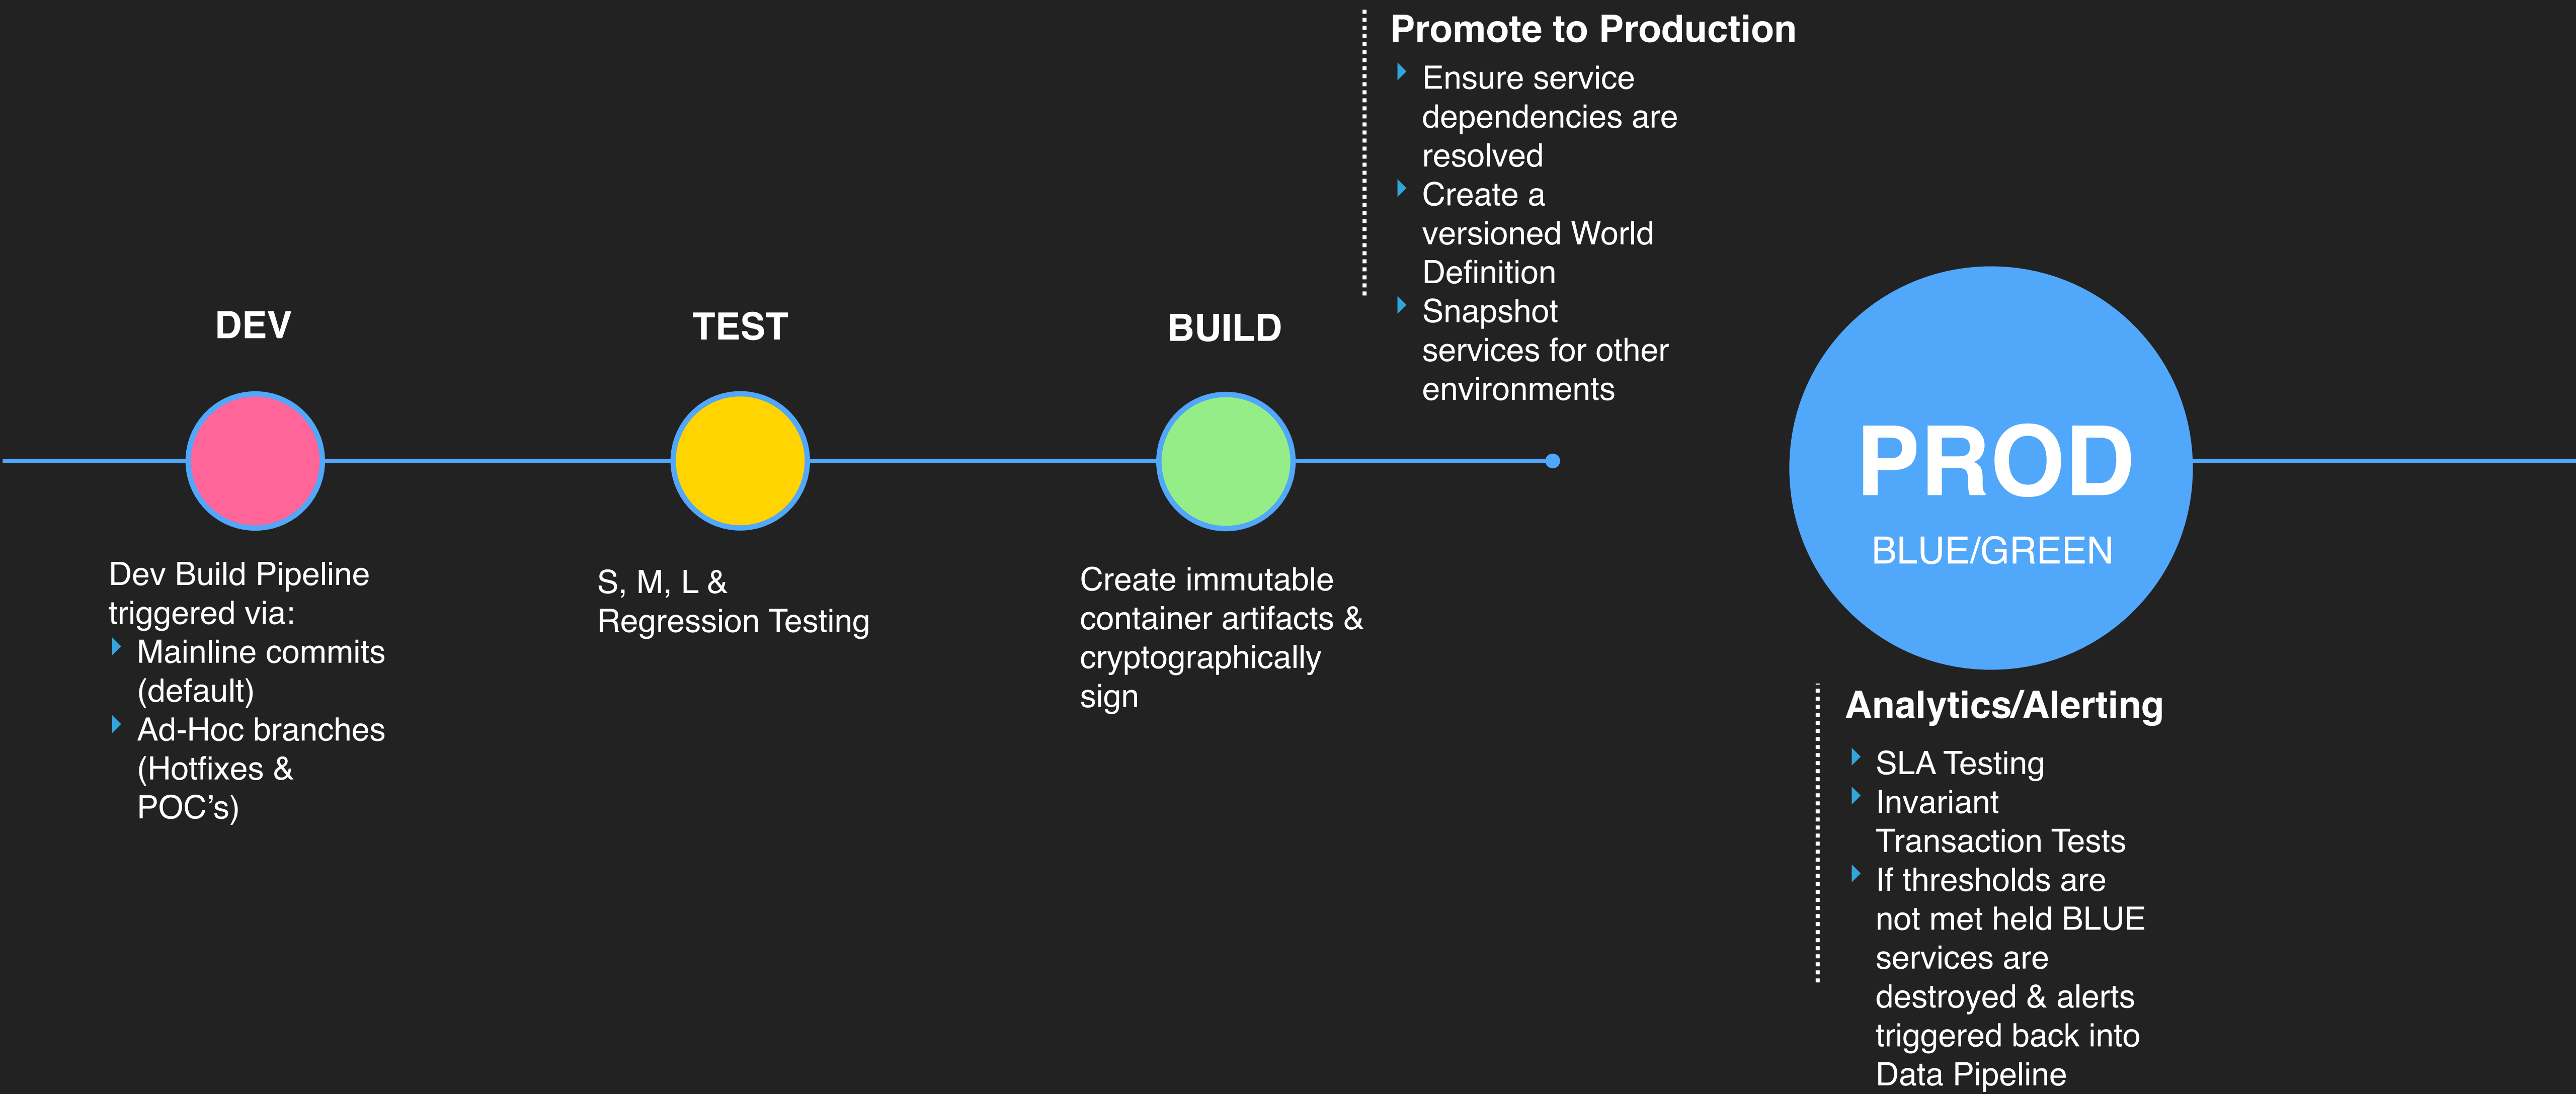


DESCRIBING OUR WORLD

- ▶ All service builds enter our “World Definition”
 - ▶ JSON definition similar to Marathon or Kubernetes catalog
 - ▶ each instance of our world is versioned for use on any environment, even local
- ▶ All orchestrated via an in-house tool we built named Tailor
 - ▶ we call it Tailor as it's our bespoke stitcher
 - ▶ created down the road in lovely Bristol, UK
 - ▶ developers issue commands such as:
`tlr world:up`



BUILD, SHIP, RUN





IN CONCLUSION

- ▶ APPLICATION SERVICES AS FULLY DISTRIBUTED ARCHITECTURE
- ▶ MICROSERVICES DON'T COME FOR FREE
- ▶ INFRASTRUCTURE IS A UNIT OF COMPUTE

www.lifion.com
jason.melo@adp.com



L I F
I O N

by

ADP®

thank you